

Wadi Moughanim - Lecture 1 (2023/10/05): Exercise 1.3

Exercise 1.3 (Domains of Attraction: Illustration)

1. Illustrate the Phenomenon of Weak Convergence

Illustrate the phenomenon of weak convergence stated in Fisher and Tippett's theorem through the convergence of histograms (built from a random sample) of maxima towards histograms of the limit, for a negative shape parameter.

- **Choose a Distribution:** Choose a textbook distribution F in the Weibull domain of attraction and find appropriate norming sequences a_n, b_n such that (MDA) holds.

We aim to show that the uniform distribution on the interval $(0,1)$ belongs to the Weibull domain of attraction. We consider the cumulative distribution function (CDF) of $U(0, 1)$,

$$F_{\text{unif}}(x) = x \cdot 1_{(0,1)}(x),$$

We proceed by choosing sequences a_n and b_n such that

$$a_n x + b_n \rightarrow \infty \quad \text{as } n \rightarrow \infty.$$

Then :

$$F^n(a_n x + b_n) = (a_n x + b_n)^n,$$

valid for all $x \leq \frac{1-b_n}{a_n}$

Let's choose $b_n = 1$ and allow a_n to be any non-zero sequence and $x < 0$:

$$F^n(a_n x + b_n) = \exp(n \log(1 + a_n x)).$$

For a specific choice of $a_n = \frac{1}{n}$, the expression simplifies to

$$F^n\left(\frac{x}{n} + 1\right) \approx \exp(x) 1_{x < 0}$$

This is a particular case of the Weibull distribution with $\alpha = -1$. Hence, with the sequences $a_n = \frac{1}{n}$ and $b_n = 1$.

-
- **Write a Short Code:** Write a short code allowing to:
 - Generate M blocks of size n of independent random variables distributed according to F and normalize the block maxima.
 - Plot a histogram of the M normalized maxima and superimpose the histogram for the limit distribution in a visually illustrative manner.

```

In [ ]: import numpy as np
import matplotlib.pyplot as plt

n = 30 # Size of each block

# Generate M blocks each of size n from U(0,1)
def generate_blocks(M, n):
    return np.random.rand(M, n)

# Compute block maxima
def compute_maxima(blocks):
    return np.max(blocks, axis=1)

# Normalize the block maxima
def normalize_maxima(maxima, n):
    a_n = 1.0 / n
    b_n = 1
    return (maxima - b_n)/a_n

def plot_analysis(blocks, block_maxima, normalized_maxima, M, nbins):
    fig, axs = plt.subplots(3, 1, figsize=(10, 9), constrained_layout=True)

    # Raw Data and Block Maxima
    block_ends = np.arange(n, M*n+1, n)
    raw_data = blocks.flatten()
    axs[0].scatter(range(len(raw_data)), raw_data, s=6, c='black', label=
    axs[0].scatter(block_ends - n/2, block_maxima, s=10, c='red', marker=
    for end in block_ends:
        axs[0].axvline(x=end, color='blue', linestyle='--')
    axs[0].set_title(f"Block Maxima and Raw Data with {M} Observations")
    axs[0].legend()

    # Normalized Maxima
    axs[1].plot(normalized_maxima, label=f"Normalized Maxima with {M} Obs
    axs[1].set_title(f"Normalized Maxima with {M} Observations")

    # Histogram and Limit Distribution
    axs[2].hist(normalized_maxima, bins=nbins, density=True, label=f"Hist
    x_vals = np.linspace(min(normalized_maxima), max(normalized_maxima),
    y_vals = np.exp(x_vals)
    axs[2].plot(x_vals, y_vals, label='Limit Distribution')
    axs[2].set_title(f"Histogram of Normalized Maxima with Limit Distribu
    axs[2].legend()

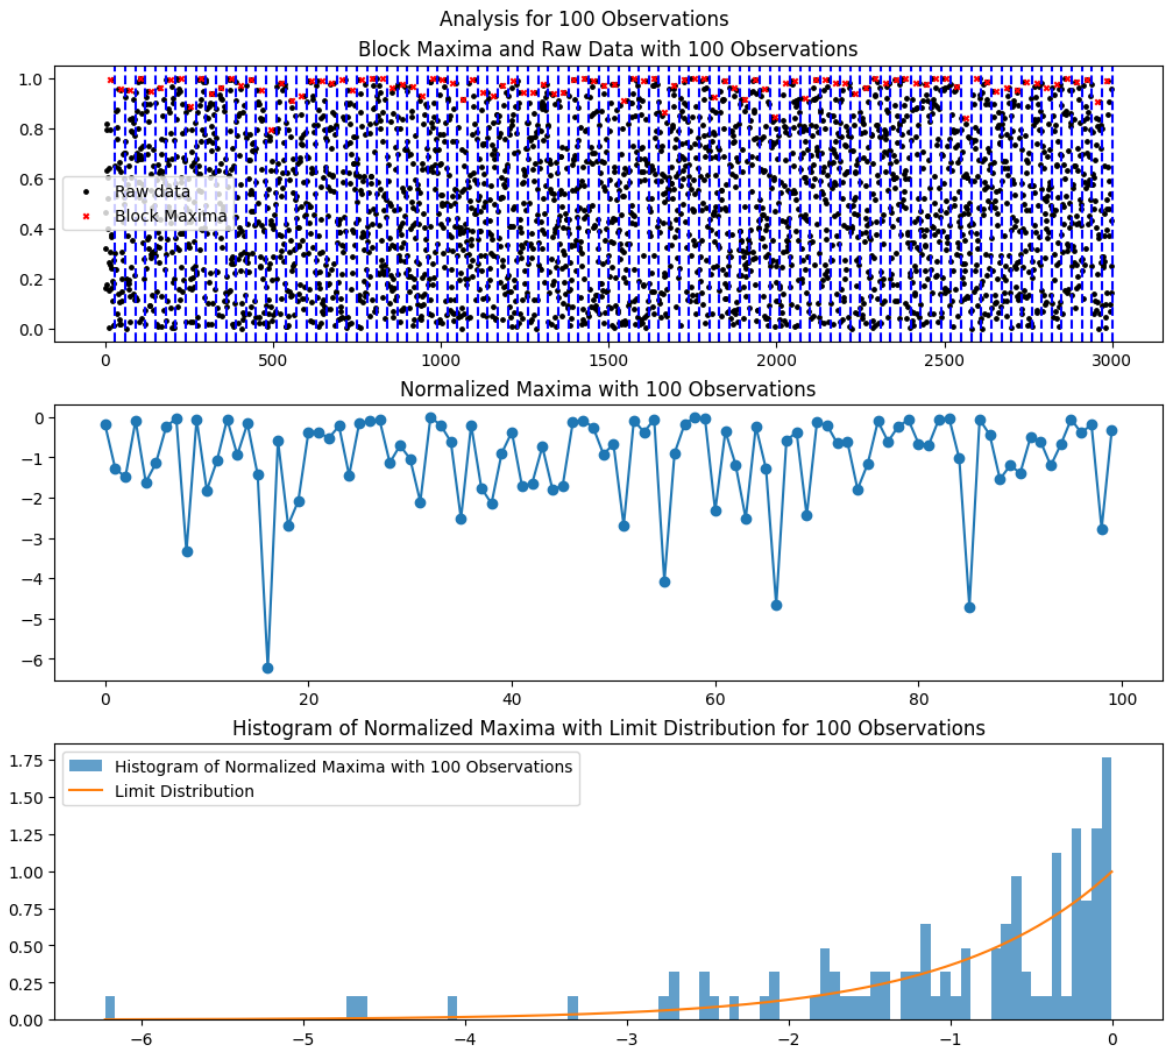
    fig.suptitle(f"Analysis for {M} Observations")
    plt.show()


# 1
M1 = 100
blocks1 = generate_blocks(M1, n)
block_maxima1 = compute_maxima(blocks1)
normalized_maxima1 = normalize_maxima(block_maxima1, n)
plot_analysis(blocks1, block_maxima1, normalized_maxima1, M1, nbins=100)

# 2
M2 = 300
blocks2 = generate_blocks(M2, n)
block_maxima2 = compute_maxima(blocks2)

```

```
normalized_maxima2 = normalize_maxima(block_maxima2, n)
plot_analysis(blocks2, block_maxima2, normalized_maxima2, M2, nbins=100)
```





- **Vary M and n:** Let M and n vary so as to illustrate weak convergence of maxima as $M \rightarrow \infty$. Explain the role of M and n in what you observe. Summarize the results in a figure including ≈ 6 such histograms with different values of M and a single (appropriate) value of n.

```
In [ ]: import numpy as np
import matplotlib.pyplot as plt

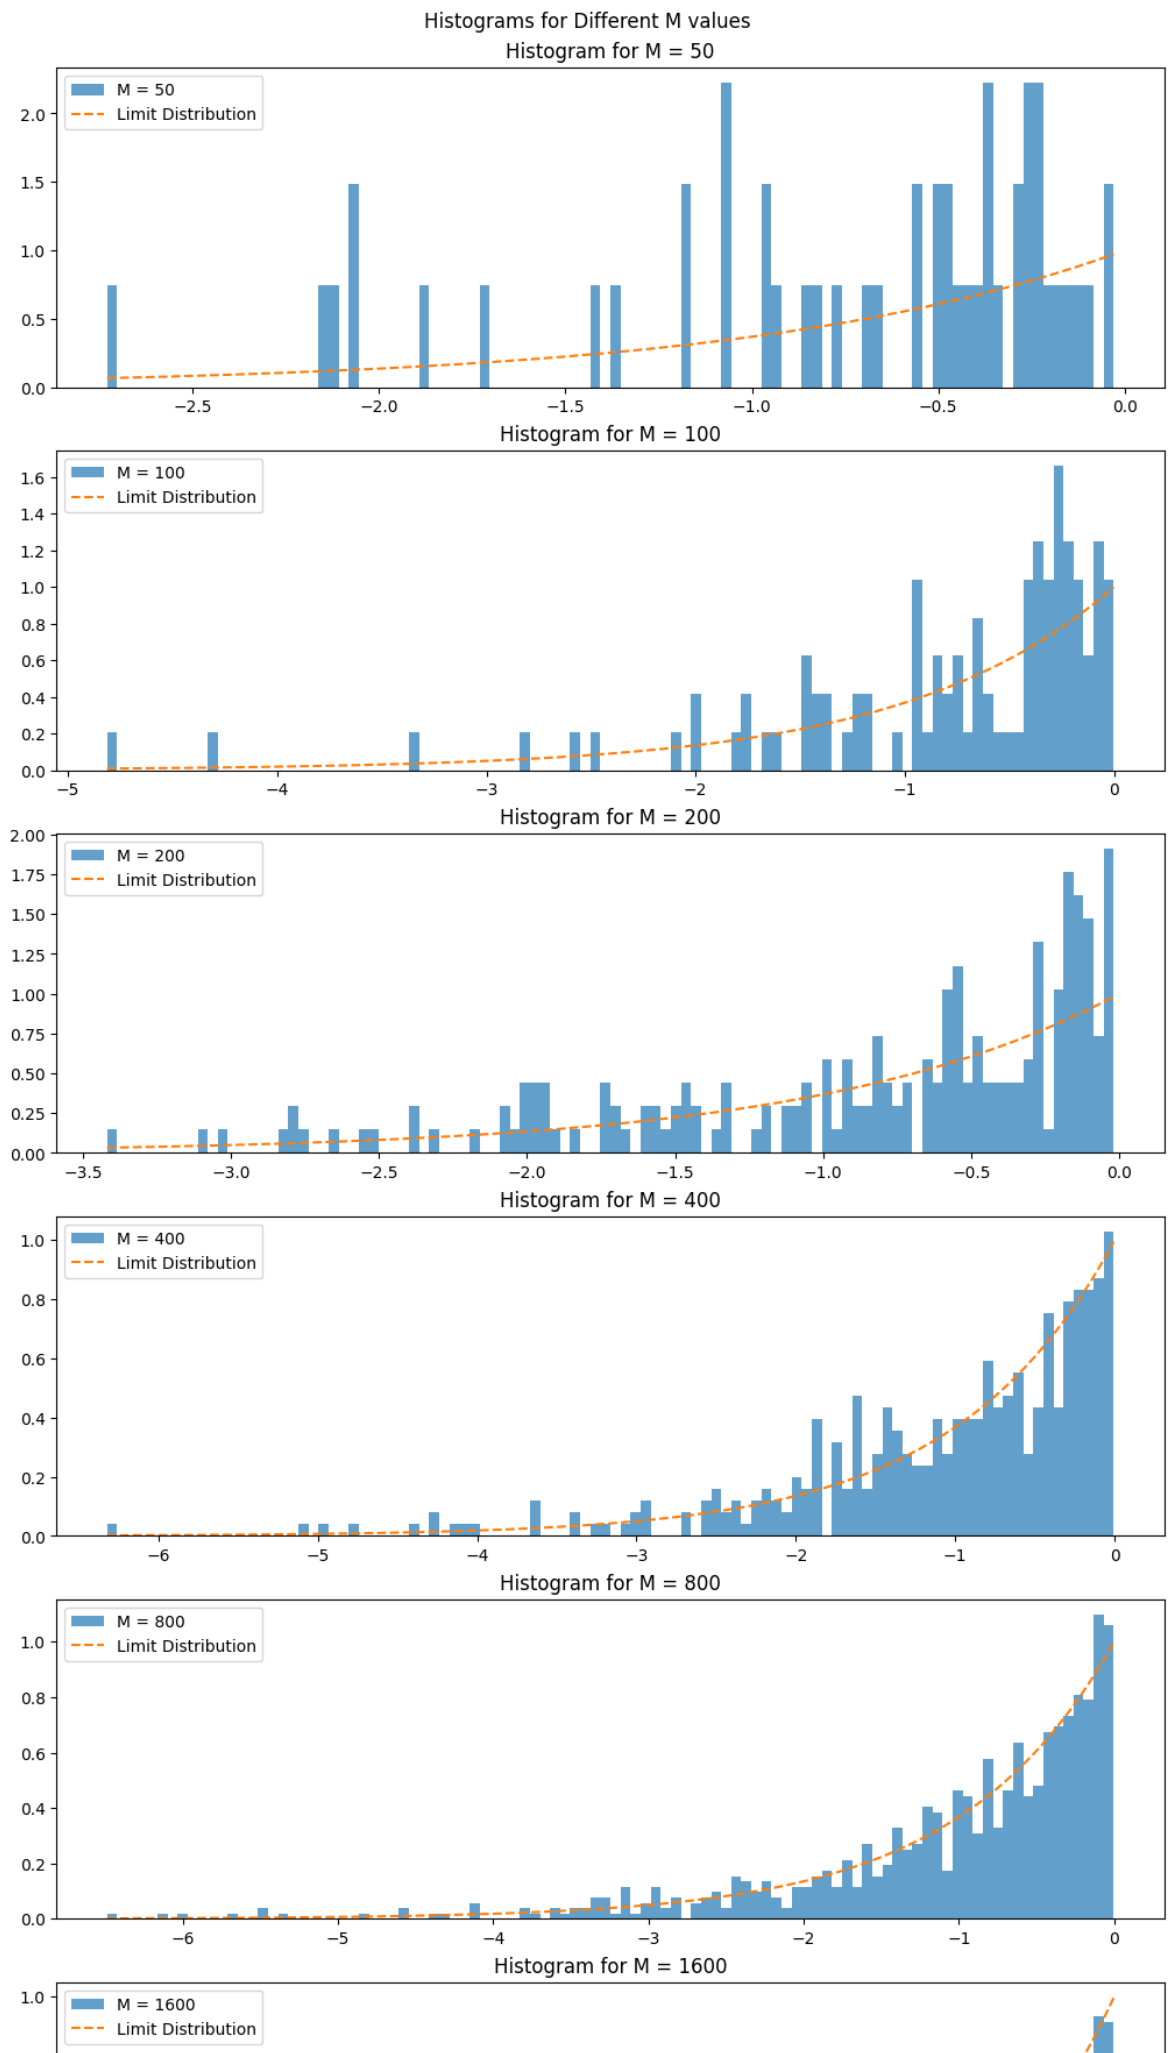
n_values = [10, 20, 40, 80, 160, 320]
M_values = [50, 100, 200, 400, 800, 1600]


# Create a new figure with 6 subplots
fig, axs = plt.subplots(len(M_values), 1, figsize=(10, 20), constrained_layout=True)

# Populate the summary figure
for idx, (M, n) in enumerate(zip(M_values, n_values)):
    blocks = generate_blocks(M, n)
    block_maxima = compute_maxima(blocks)
    normalized_maxima = normalize_maxima(block_maxima, n)
    x_vals = np.linspace(min(normalized_maxima), max(normalized_maxima), 100)
    y_vals = np.exp(x_vals)

    # Histogram
    axs[idx].hist(normalized_maxima, bins=100, density=True, label=f'M = {M}')
    # Limit Distribution
    axs[idx].plot(x_vals, y_vals, label='Limit Distribution', linestyle='solid')
```

```
    axs[idx].set_title(f"Histogram for M = {M}")  
    axs[idx].legend()  
  
fig.suptitle("Histograms for Different M values")  
plt.show()
```





2. Uniform Convergence of c.d.f's

Show (graphically and numerically) uniform convergence of c.d.f's. Explain why (i.e., prove that) weak convergence of normalized maxima indeed implies uniform convergence of c.d.f's.

The cdf. F_n of the normalized maxima converges to F at all cadlag points of F . Given that F is cadlag everywhere in the 3 cases Fréchet, Weibull and Gumbel, F_n converges to F for every $x \in \mathbb{R}$.

As F is increasing and for any $\epsilon > 0$, we take points

$-\infty = x_0 < x_1 < \dots < x_k = +\infty$ such that $\forall i \in \{0, \dots, k-1\}$

$F(x_{i+1}) - F(x_i) \leq \epsilon$

For any sub-interval $[x_i, x_{i+1}]$ of $[x_0, x_k]$, we have $F(x_{i+1}) - F(x_i) \leq \epsilon$. Thus as both F_n and F are increasing and bounded, we have

$\sup_{x \in [x_i, x_{i+1}]} |F_n(x) - F(x)| \leq \max(|F_n(x_i) - F(x_i)|, |F_n(x_{i+1}) - F(x_{i+1})|) + \epsilon$

.

$$\text{Thus : } \sup_{x \in \mathbb{R}} |F_n(x) - F(x)| \leq \max_{j=0,1,\dots,K} |F_n(x_j) - F(x_j)| + \epsilon$$

As $n \rightarrow \infty$, the term $\max_{j=0,1,\dots,K} |F_n(x_j) - F(x_j)|$ goes to 0 due to weak convergence. Thus, for sufficiently large n , the supremum of the absolute differences can be made smaller than any given ϵ , implying uniform convergence.

```
In [ ]: import numpy as np
import matplotlib.pyplot as plt

# Parameters
n = 30 # Size of each block

def generate_blocks(M, n):
    return np.random.rand(M, n)

def compute_maxima(blocks):
    return np.max(blocks, axis=1)

def normalize_maxima(maxima, n):
    a_n = 1.0 / n
    b_n = 1
    return (maxima - b_n) / a_n

def empirical_cdf(data):
    sorted_data = np.sort(data)
```

```

y = np.arange(1, len(data) + 1) / len(data)
return sorted_data, y

def weibull_cdf(x):
    return np.exp(x)

def plot_convergence(M_values):
    fig, axs = plt.subplots(len(M_values), 1, figsize=(8, 4*len(M_values)))

    for i, M in enumerate(M_values):
        blocks = generate_blocks(M, n)
        maxima = compute_maxima(blocks)
        normalized_maxima = normalize_maxima(maxima, n)

        x_empirical, y_empirical = empirical_cdf(normalized_maxima)

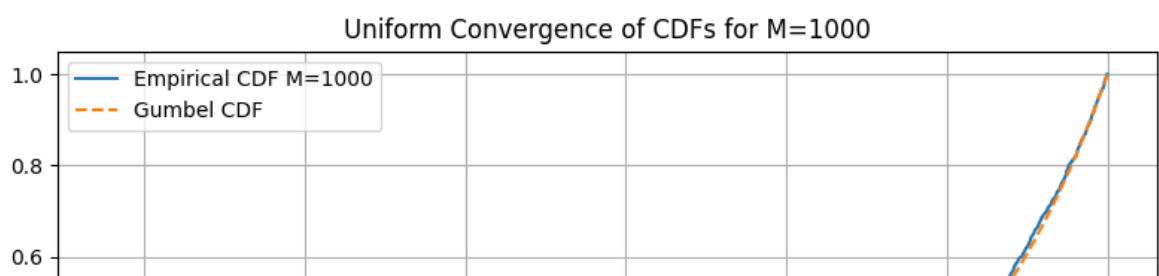



        x_vals = np.linspace(min(normalized_maxima), max(normalized_maxima), 100)
        y_vals = weibull_cdf(x_vals)


        axs[i].plot(x_empirical, y_empirical, label=f'Empirical CDF M={M}')
        axs[i].plot(x_vals, y_vals, '--', label='Gumbel CDF')
        axs[i].legend()
        axs[i].grid(True)
        axs[i].set_title(f'Uniform Convergence of CDFs for M={M}')

    plt.tight_layout()
    plt.show()


M_values = [30, 100, 300, 1000, 3000, 10000]
plot_convergence(M_values)

```







Uniform Convergence of CDFs for M=3000



Uniform Convergence of CDFs for M=10000



3. Translated Pareto Distribution

Change the input distribution and work with a translated Pareto distribution,

$P(X > x) = \left(\frac{x-\beta}{u}\right)^{-\alpha}$, for some $\alpha > 0, \beta, u \in \mathbb{R}, x \geq u + \beta$. Draw similar outputs as in the previous questions and compare the rate of convergence.

$$F(x) = 1 - \left(\frac{x-\beta}{u}\right)^{-\alpha} \quad \text{for } x > u + \beta,$$

we are looking to find sequences a_n and b_n that normalize maxima.

$$P(X > x) = \left(\frac{x - \beta}{u} \right)^{-\alpha}.$$

Given the normalization $F_n(a_n x + b_n)$, and setting $b_n = \beta$:

$$F_n(a_n x + \beta) = \left(1 - \left(\frac{a_n x + \beta - \beta}{u} \right)^{-\alpha} \right)^n$$

$$F_n(a_n x + \beta) = \left(1 - \left(\frac{a_n x}{u} \right)^{-\alpha} \right)^n.$$

For large n , this behaves as:

$$F_n(a_n x + \beta) \approx \exp \left(n \log \left(1 - \left(\frac{a_n x}{u} \right)^{-\alpha} \right) \right).$$

For the normalization to converge to the Fréchet distribution, we aim for a structure that's proportional to $\exp(-x^{-\alpha})$.

To attain this, let's choose $a_n = un^{1/\alpha}$. Plugging in this value:

$$F_n(un^{1/\alpha}x + \beta) \approx \exp \left(n \log \left(1 - \left(\frac{n^{1/\alpha}x}{u} \right)^{-\alpha} \right) \right).$$

For large n , the term inside the logarithm gives us the desired structure, $\exp(-x^{-\alpha})$, which is the Fréchet distribution.

```
In [ ]: import numpy as np
import matplotlib.pyplot as plt

# Parameters
n = 30 # Size of each block

# List of parameter sets
parameter_sets = [
    {"alpha": 2.0, "beta": 0.5, "u": 1.0},
    {"alpha": 2.5, "beta": 0.7, "u": 1.2},
    {"alpha": 3.0, "beta": 0.9, "u": 1.5}
]

def generate_blocks(M, n, alpha, beta, u):
    uniform_samples = np.random.rand(M, n)
    pareto_samples = beta + u * ((1 - uniform_samples)**(-1/alpha))
    return pareto_samples

def compute_maxima(blocks):
    return np.max(blocks, axis=1)

def normalize_maxima(maxima, n):
    a_n = u * n**(1/alpha)
    b_n = beta
    return (maxima - b_n) / a_n

def plot_analysis(blocks, block_maxima, normalized_maxima, M, nbins, alpha):
    block_ends = np.arange(n, M*n+1, n)
```

```

raw_data = blocks.flatten()

fig, axs = plt.subplots(3, 1, figsize=(10, 15))

axs[0].plot(raw_data, 'ko', markersize=6, label='Raw data')
axs[0].plot(block_ends-n/2, block_maxima, 'rx', markersize=10, label='Block Maxima')
for end in block_ends:
    axs[0].axvline(x=end, color='blue', linestyle='--')
axs[0].legend()
axs[0].set_title(f"Block Maxima and Raw Data with {M} Observations")

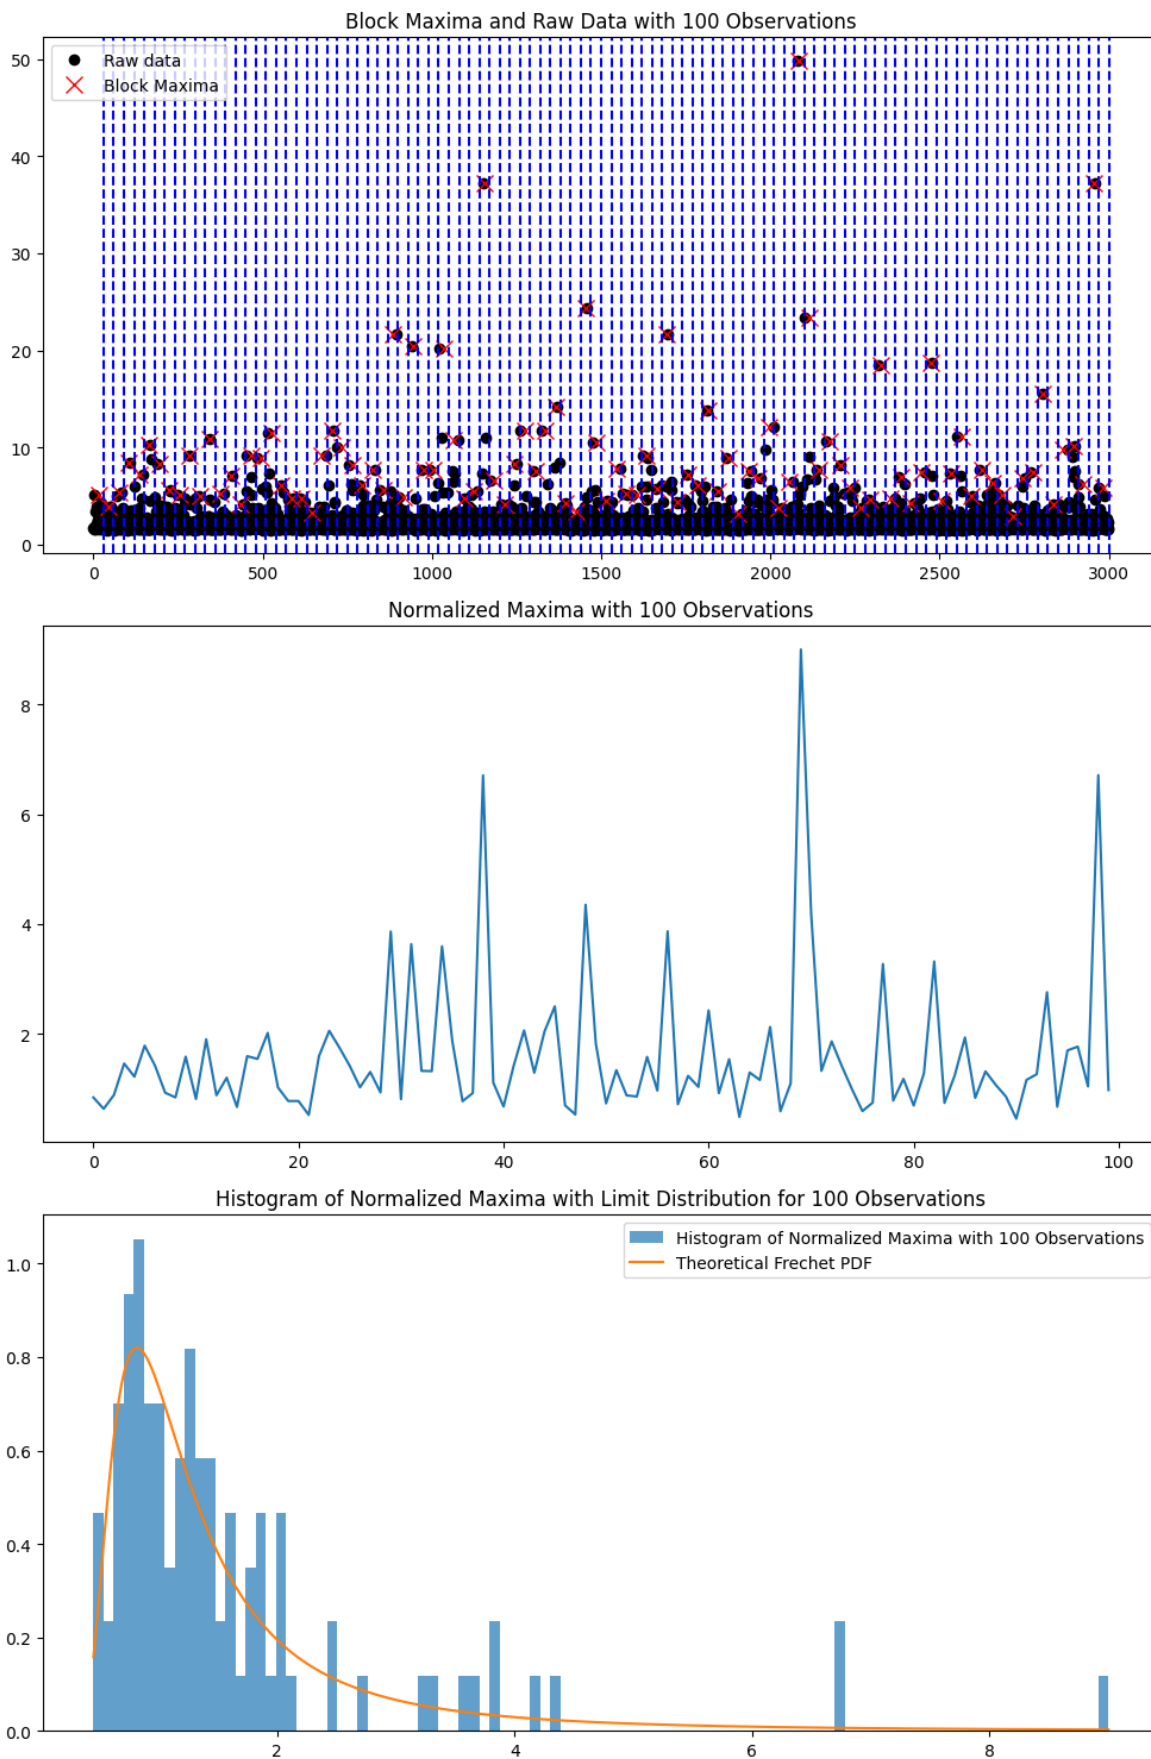
axs[1].plot(normalized_maxima, label=f"Normalized Maxima with {M} Observations")
axs[1].set_title(f"Normalized Maxima with {M} Observations")

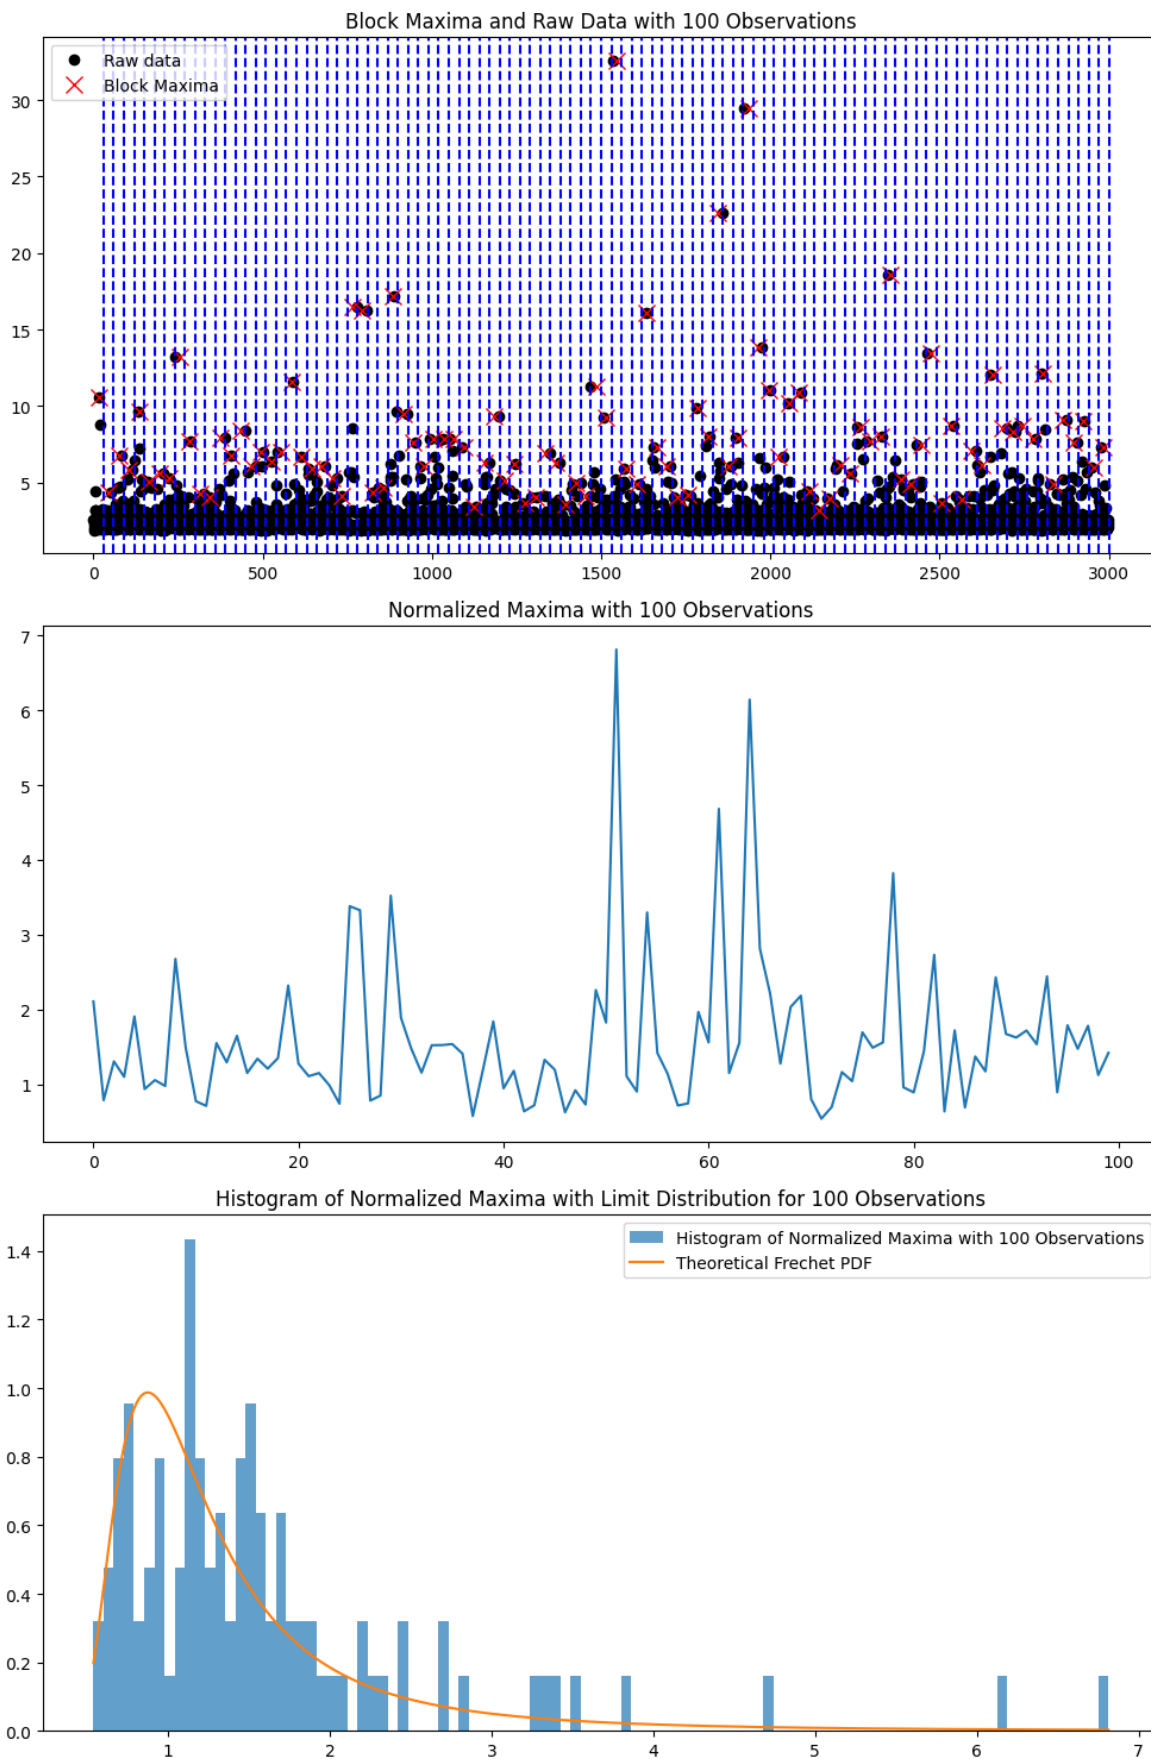
axs[2].hist(normalized_maxima, bins=nbins, density=True, label=f"Histogram of Normalized Maxima with Limit Distribution")
x_vals = np.linspace(min(normalized_maxima), max(normalized_maxima), nbins)
y_vals = alpha * x_vals**(-alpha - 1) * np.exp(-x_vals**(-alpha))
axs[2].plot(x_vals, y_vals, label='Theoretical Frechet PDF')
axs[2].set_title(f"Histogram of Normalized Maxima with Limit Distribution and Theoretical Frechet PDF")
axs[2].legend()

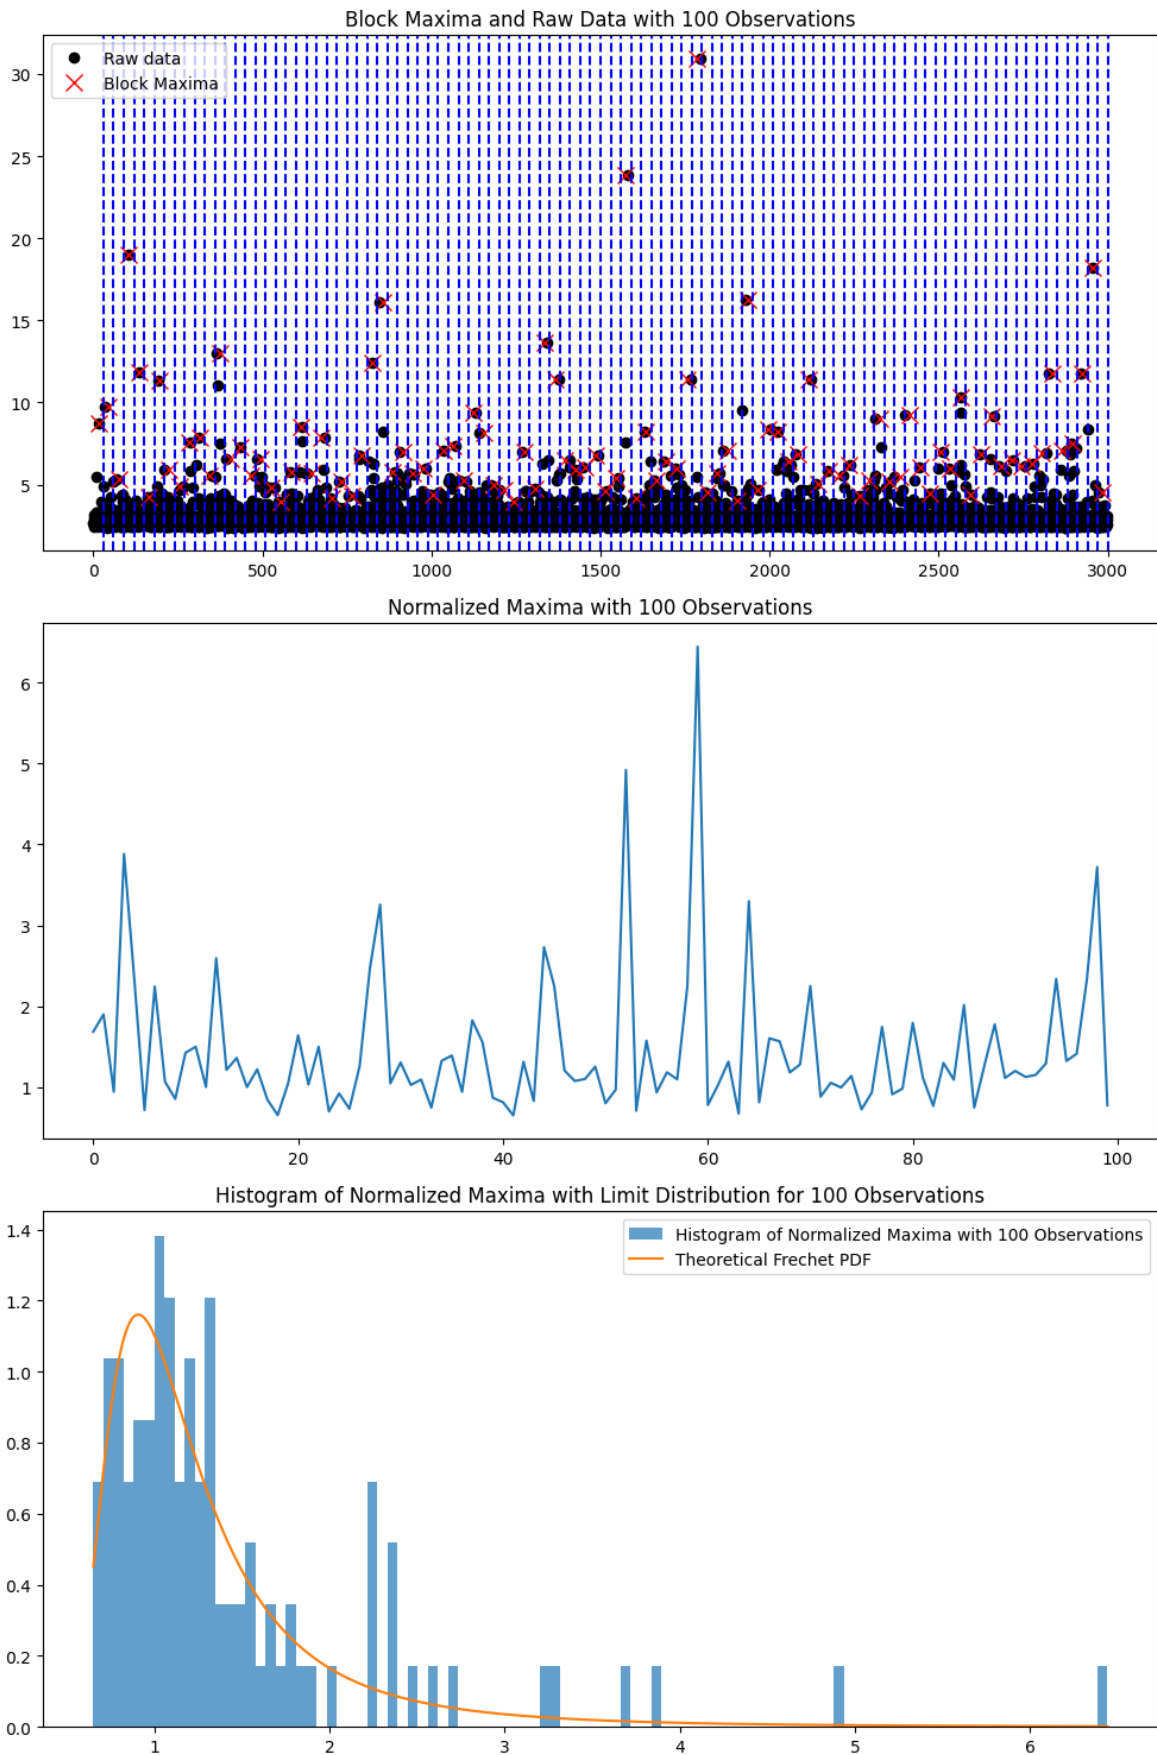
plt.tight_layout()
plt.show()

for params in parameter_sets:
    alpha, beta, u = params["alpha"], params["beta"], params["u"]
    blocks = generate_blocks(M1, n, alpha, beta, u)
    block_maxima = compute_maxima(blocks)
    normalized_maxima = normalize_maxima(block_maxima, n)
    plot_analysis(blocks, block_maxima, normalized_maxima, M1, nbins=100,

```







```
In [ ]: import numpy as np
import matplotlib.pyplot as plt

# Parameters
n = 30 # Size of each block

# List of parameter sets
```

```

parameter_sets = [
    {"alpha": 2.0, "beta": 0.5, "u": 1.0},
    {"alpha": 2.5, "beta": 0.7, "u": 1.2},
    {"alpha": 3.0, "beta": 0.9, "u": 1.5}
]

def generate_blocks(M, n, alpha, beta, u):
    uniform_samples = np.random.rand(M, n)
    pareto_samples = beta + u * ((1 - uniform_samples)**(-1/alpha))
    return pareto_samples

def compute_maxima(blocks):
    return np.max(blocks, axis=1)

def normalize_maxima(maxima, n, alpha, u, beta):
    a_n = u * n**(1/alpha)
    b_n = beta
    return (maxima - b_n) / a_n

def empirical_cdf(data):
    sorted_data = np.sort(data)
    y = np.arange(1, len(data) + 1) / len(data)
    return sorted_data, y

def frechet_cdf(x, alpha):
    return np.exp(-x**(-alpha))

def plot_convergence(M_values, alpha, beta, u):
    fig, axs = plt.subplots(len(M_values), 1, figsize=(8, 4*len(M_values)))

    for i, M in enumerate(M_values):
        blocks = generate_blocks(M, n, alpha, beta, u)
        maxima = compute_maxima(blocks)
        normalized_maxima = normalize_maxima(maxima, n, alpha, u, beta)

        x_empirical, y_empirical = empirical_cdf(normalized_maxima)

        x_vals = np.linspace(min(normalized_maxima), max(normalized_maxima))
        y_vals = frechet_cdf(x_vals, alpha)

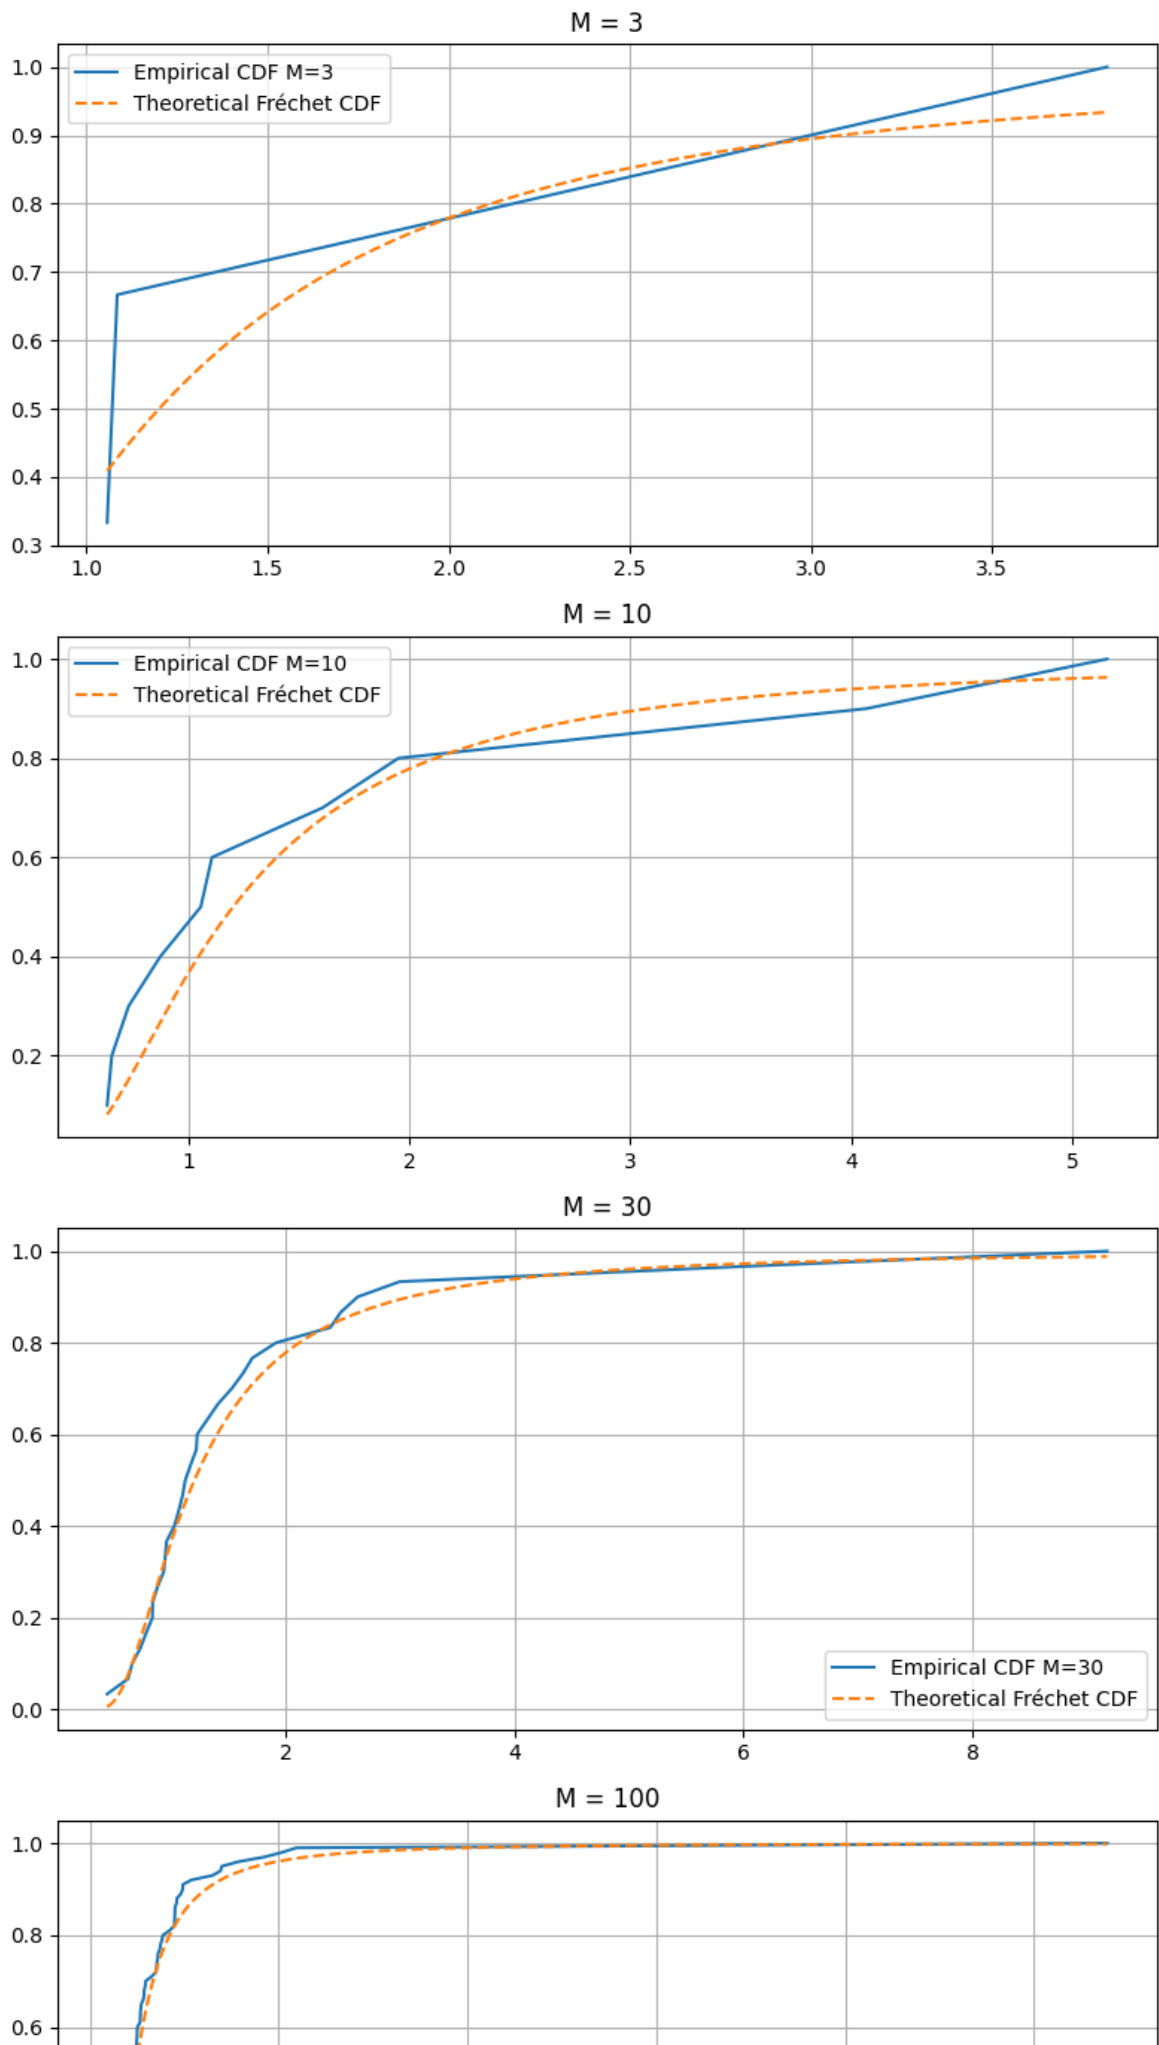
        axs[i].plot(x_empirical, y_empirical, label=f'Empirical CDF M={M}')
        axs[i].plot(x_vals, y_vals, '--', label='Theoretical Fréchet CDF')
        axs[i].legend()
        axs[i].set_title(f'M = {M}')
        axs[i].grid(True)


    plt.tight_layout()
    plt.show()

# List of sample sizes to analyze
M_values = [3, 10, 30, 100, 300, 1000]


# Iterate over parameter sets and plot the convergence for each
for params in parameter_sets:
    alpha, beta, u = params["alpha"], params["beta"], params["u"]
    plot_convergence(M_values, alpha, beta, u)

```









$M = 300$




$M = 1000$










$M = 300$




$M = 1000$








M = 300



M = 1000



Now let's see the difference between the rates of convergence

```
In [ ]: import matplotlib.pyplot as plt

def ks_statistic(empirical_x, empirical_y, theoretical_cdf, alpha):
    """Compute the Kolmogorov-Smirnov statistic."""
    theoretical_y = [theoretical_cdf(x, alpha) for x in empirical_x]
    return np.max(np.abs(empirical_y - theoretical_y))

def plot_ks_statistics(M_values, parameter_sets):
    """Plot the KS statistics for different sample sizes and parameter sets"""
    plt.figure(figsize=(10, 6))

    for params in parameter_sets:
```

```

alpha, beta, u = params["alpha"], params["beta"], params["u"]
ks_stats = []

for M in M_values:
    blocks = generate_blocks(M, n, alpha, beta, u)
    maxima = compute_maxima(blocks)
    normalized_maxima = normalize_maxima(maxima, n, alpha, u, beta)
    x_empirical, y_empirical = empirical_cdf(normalized_maxima)


    ks = ks_statistic(x_empirical, y_empirical, frechet_cdf, alpha)
    ks_stats.append(ks)

plt.plot(M_values, ks_stats, '-o', label=f' $\alpha$ ={alpha},  $\beta$ ={beta}, u={u}')

plt.xscale("log")
plt.title('Rate of Convergence')
plt.xlabel('Sample Size')
plt.ylabel('KS Statistic')
plt.legend()
plt.grid(True, which="both", ls="--")
plt.tight_layout()
plt.show()

plot_ks_statistics(M_values, parameter_sets)

```



4. Estimate the GEV Parameters

With the distribution from question 1 or 3, generate a dataset of an appropriate size and estimate the GEV parameters with a maximum-likelihood method. Discuss the convergence towards the true parameters.

Let's pick up the pareto distribution from question 3 and generate a dataset of size 10000.

```

In [ ]: M = 10000
        alpha_true = 2.0

```

```

beta_true = 0.5
u_true = 1.0

blocks = generate_blocks(M, n, alpha_true, beta_true, u_true)
maxima = compute_maxima(blocks)

from scipy.stats import genextreme as gev
from scipy.optimize import minimize

# Negative log-likelihood for the GEV distribution
def negative_gev_log_likelihood(params, data):
    c, loc, scale = params
    return -np.sum(gev.logpdf(data, c, loc=loc, scale=scale))

# Maximum-likelihood estimation
initial_guess = [0.5, np.mean(maxima), np.std(maxima)]
res = minimize(negative_gev_log_likelihood, initial_guess, args=(maxima,))
c_hat, loc_hat, scale_hat = res.x

print(f"Estimated parameters: c = {c_hat:.4f}, loc = {loc_hat:.4f}, scale = {scale_hat:.4f}")

```

Estimated parameters: c = 0.5000, loc = 10.4209, scale = 15.9310

/Users/wadimoughanim/Library/Python/3.9/lib/python/site-packages/scipy/optimize/_numdiff.py:576: RuntimeWarning:

invalid value encountered in subtract

We obtain a shape value of 0.5000, a location (which is a bit like an average) of 10.3723 and a scale (which indicates how scattered our data are) of 15.1434.